

Introduction

The Wasserstein Generative Adversarial Network (WGAN) introduces a new approach to training GANs, aiming to improve stability and address vanishing gradients by minimizing the Wasserstein distance between the data and generated distributions.

Key Equations

The Wasserstein distance W can be expressed as:

$$W(P_r, P_g) = \inf_{\gamma \in \Pi(P_r, P_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

where P_r is the real data distribution, P_g is the generated data distribution, and $\Pi(P_r, P_g)$ is the set of all joint distributions with marginals P_r and P_g .

Algorithm 1: WGAN Training Procedure

Detailed Explanation

Critic Update (Lines 3–7)

- Sample a batch of real data points:

$$\{x^{(i)}\}_{i=1}^m \sim P_r.$$

- Sample a batch of generated data points:

$$z^{(i)} \sim P_z, \quad \{g_\theta(z^{(i)})\}_{i=1}^m.$$

- Compute the gradient of the critic loss:

$$g_w = \nabla_w \left[\frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right].$$

- Update the critic's parameters:

$$w \leftarrow w + \alpha \cdot \text{RMSProp}(w, g_w).$$

- Clip the critic's weights:

$$w \leftarrow \text{clip}(w, -c, c).$$

- Repeat the critic update n_{critic} times: This ensures the critic sufficiently approximates the Wasserstein distance before updating the generator.

Algorithm 1 WGAN Algorithm

Require: α : learning rate, c : clipping parameter, m : batch size, n_{critic} : number of critic updates per generator update.

Require: w_0 : initial critic parameters, θ_0 : initial generator parameters.

- 1: **while** θ has not converged **do**
- 2: **for** $t = 0, \dots, n_{\text{critic}} - 1$ **do**
- 3: Sample $\{x^{(i)}\}_{i=1}^m \sim P_r$: a batch of real data points.
- 4: Sample $\{z^{(i)}\}_{i=1}^m \sim P_z$: a batch of prior samples.
- 5: Compute the gradient of the critic loss:

$$g_w = \nabla_w \left[\frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right].$$

- 6: Update critic parameters:

$$w \leftarrow w + \alpha \cdot \text{RMSPProp}(w, g_w).$$

- 7: Clip critic weights:

$$w \leftarrow \text{clip}(w, -c, c).$$

- 8: **end for**
- 9: Sample $\{z^{(i)}\}_{i=1}^m \sim P_z$: a batch of prior samples.
- 10: Compute the gradient of the generator loss:

$$g_\theta = -\nabla_\theta \left[\frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right].$$

- 11: Update generator parameters:

$$\theta \leftarrow \theta - \alpha \cdot \text{RMSPProp}(\theta, g_\theta).$$

- 12: **end while**
-

Generator Update (Lines 8–11)

- Sample a batch of prior samples:

$$\{z^{(i)}\}_{i=1}^m \sim P_z, \quad \{g_\theta(z^{(i)})\}_{i=1}^m.$$

- Compute the gradient of the generator loss:

$$g_\theta = -\nabla_\theta \left[\frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right].$$

- Update the generator's parameters:

$$\theta \leftarrow \theta - \alpha \cdot \text{RMSPProp}(\theta, g_\theta).$$

Key Concepts and Math

Wasserstein Distance

The WGAN minimizes the Wasserstein-1 distance between the real distribution P_r and the generated distribution P_g :

$$W(P_r, P_g) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim P_r}[f(x)] - \mathbb{E}_{x \sim P_g}[f(x)],$$

where $\|f\|_L \leq 1$ enforces that f (the critic) is 1-Lipschitz. The clipping of w approximates this constraint.

Critic Loss

The critic aims to maximize the difference between real and generated scores:

$$L_w = \mathbb{E}_{x \sim P_r}[f_w(x)] - \mathbb{E}_{z \sim P_z}[f_w(g_\theta(z))].$$

Generator Loss

The generator minimizes the critic's score for generated samples:

$$L_\theta = -\mathbb{E}_{z \sim P_z}[f_w(g_\theta(z))].$$

Optimization

The algorithm uses RMSPProp for optimization. This adjusts the step size for each parameter based on the magnitude of recent gradients:

$$\text{RMSPProp}(x, g) = \frac{g}{\sqrt{\mathbb{E}[g^2]} + \epsilon},$$

where ϵ is a small constant to prevent division by zero.

Clipping

The critic's weights are clipped to lie within $[-c, c]$ element-wise to ensure Lipschitz continuity:

$$\text{clip}(w, -c, c) = \max(-c, \min(c, w)).$$

Termination Condition

The algorithm continues until the generator parameters θ converge, i.e., when the generator produces high-quality samples indistinguishable from the real distribution.